

Music generation with GAN and RL

Ivan Anokhin, Egor Nuzhin
2018



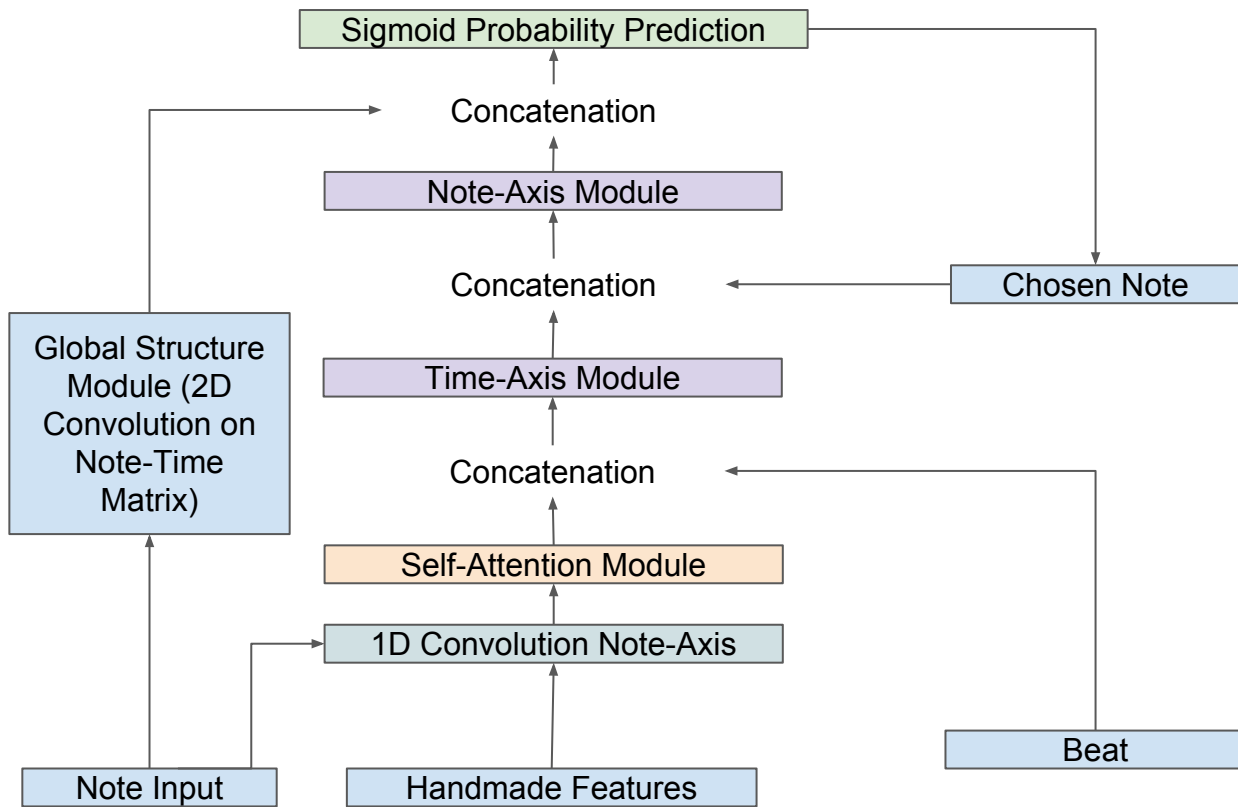
Introduction

- Listening to music is one of important part of our life
- Music is a simplest way to relax
- Music industry is one of the biggest market place it the world
- Musical composition are required in many different areas: cinema, games, ets.

Dataset description

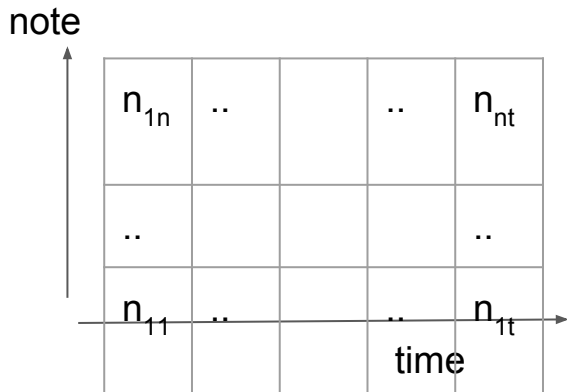
- We used open source dataset: The Lakh MIDI Dataset
- Transformed structure of midi file to use in model is:
 - Play matrix
 - Replay matrix
 - Volume matrix

Generator Architecture



Note input and Convolution Modules

1. Note Input - Note by Time Matrix with note embedding in each position



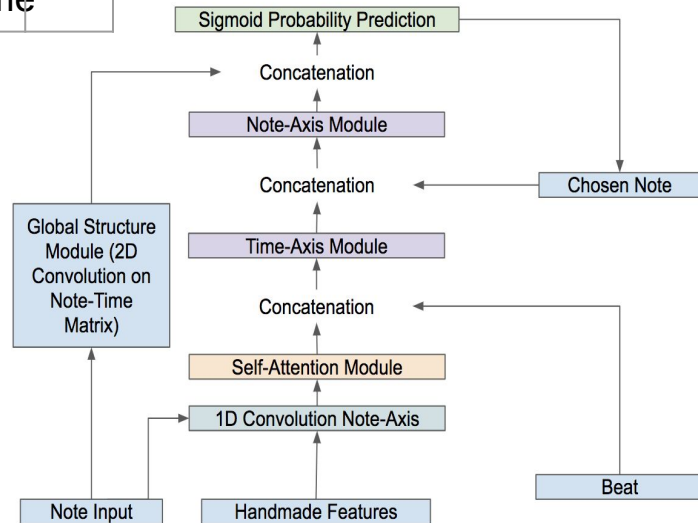
2. 1D Convolution Note-Axis (Kernel size = 24)

1D Convolution Note-Axis

3. Global Structure Module (Kernel size = (all note x 32 time steps), stride = 16)

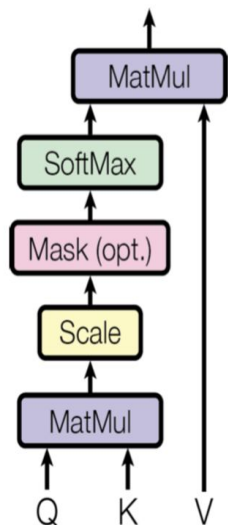
Global Structure Module (2D Convolution on Note-Time Matrix)

Note Input



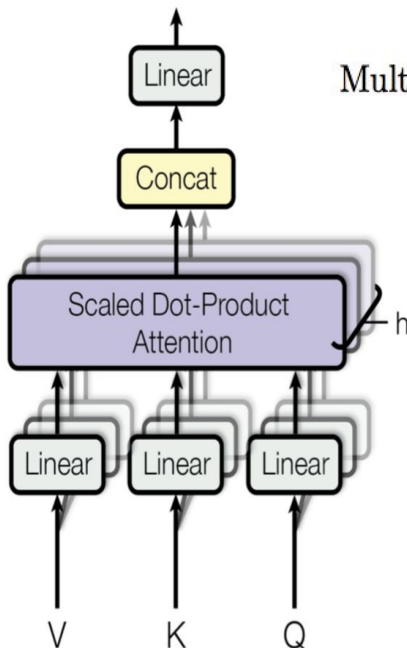
Self-Attention module

Scaled Dot-Product Attention



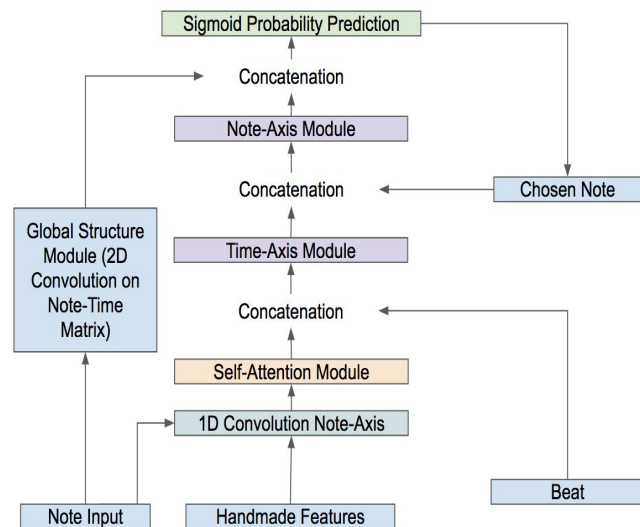
Self Attention Mechanism: When $Q=K=V$

Multi-Head Attention

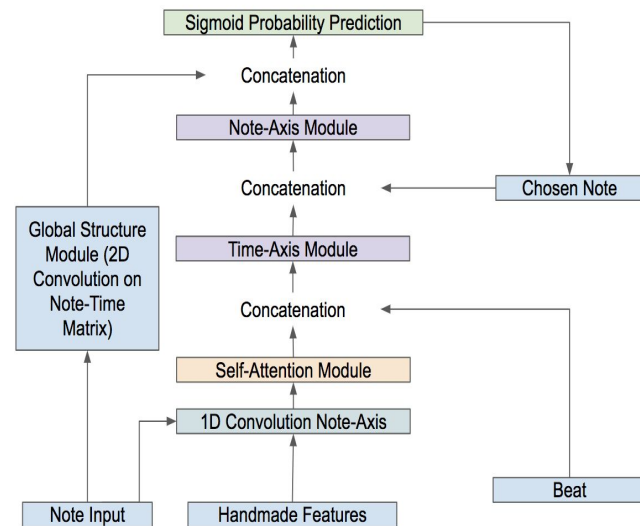
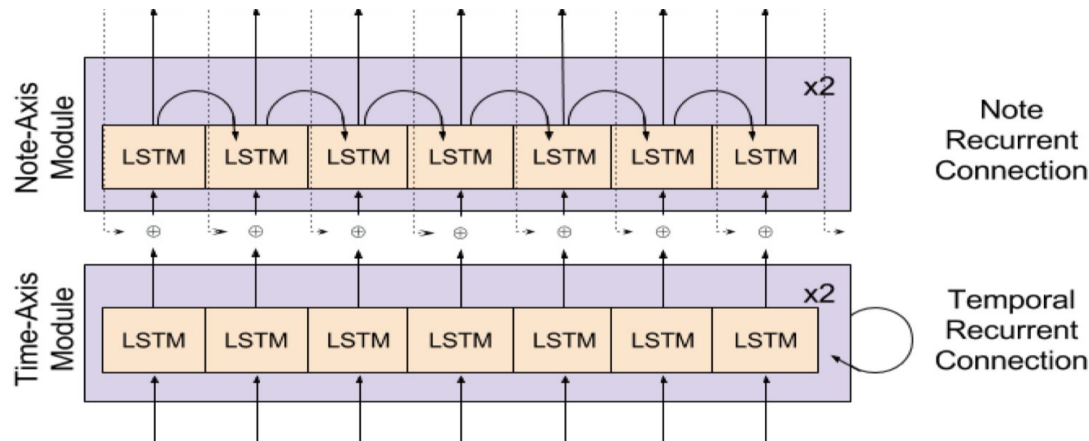


$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O$$

where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$

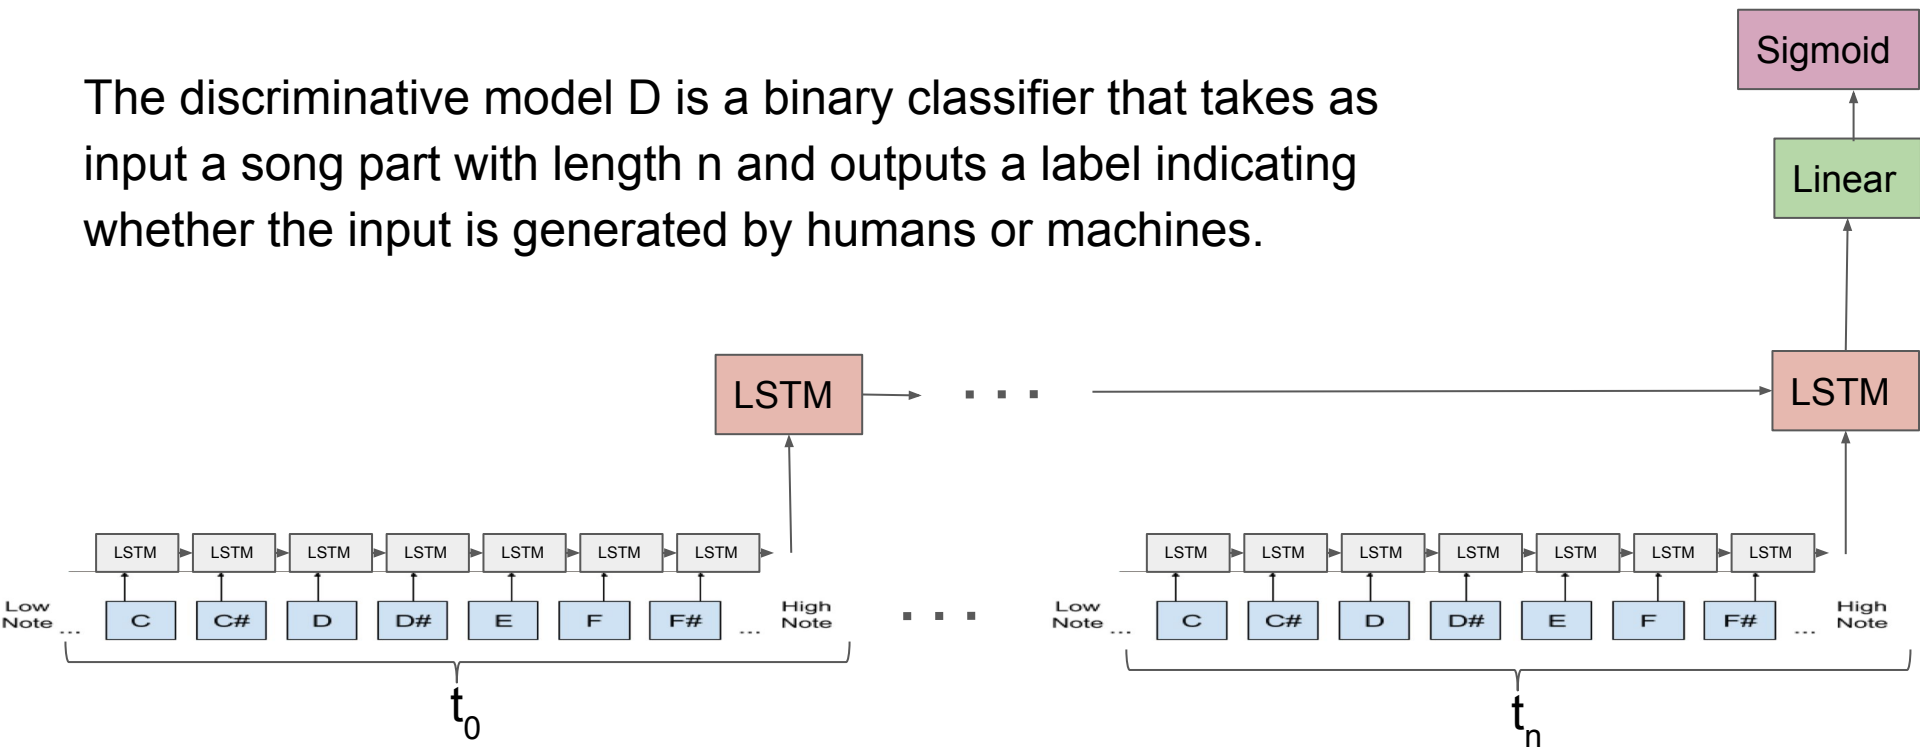


Time-Axis and Note-Axis Modules



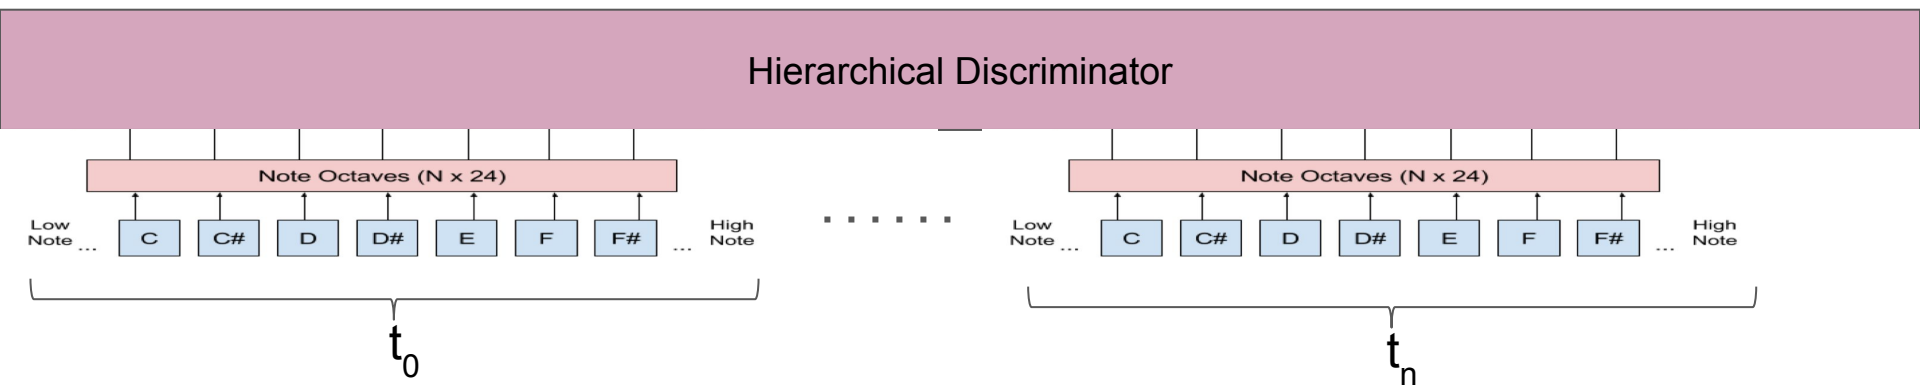
Discriminator hierarchical architecture

The discriminative model D is a binary classifier that takes as input a song part with length n and outputs a label indicating whether the input is generated by humans or machines.

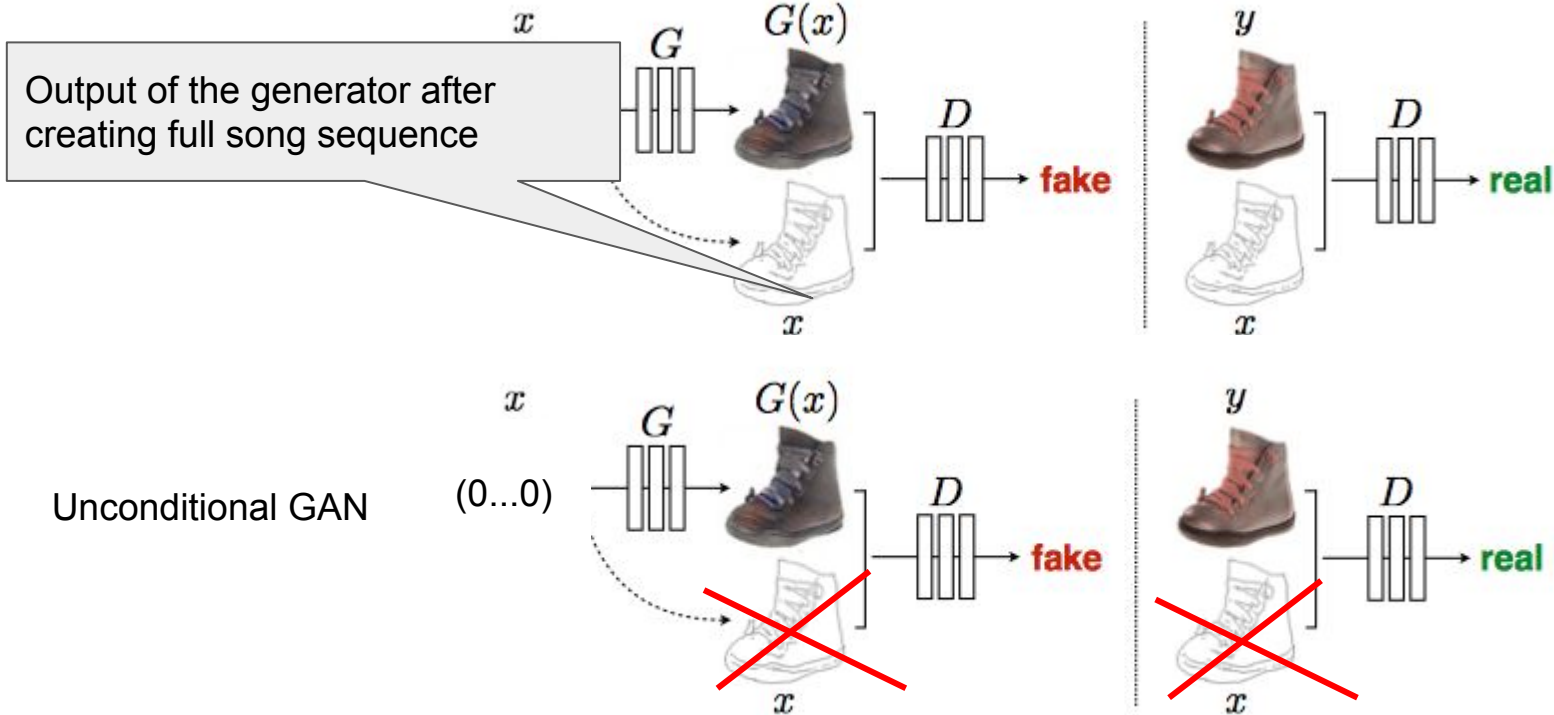


Discriminator hierarchical architecture

As additional improvement we used Note Octave embedding



Discriminator inputs



Learning technique. Algorithm

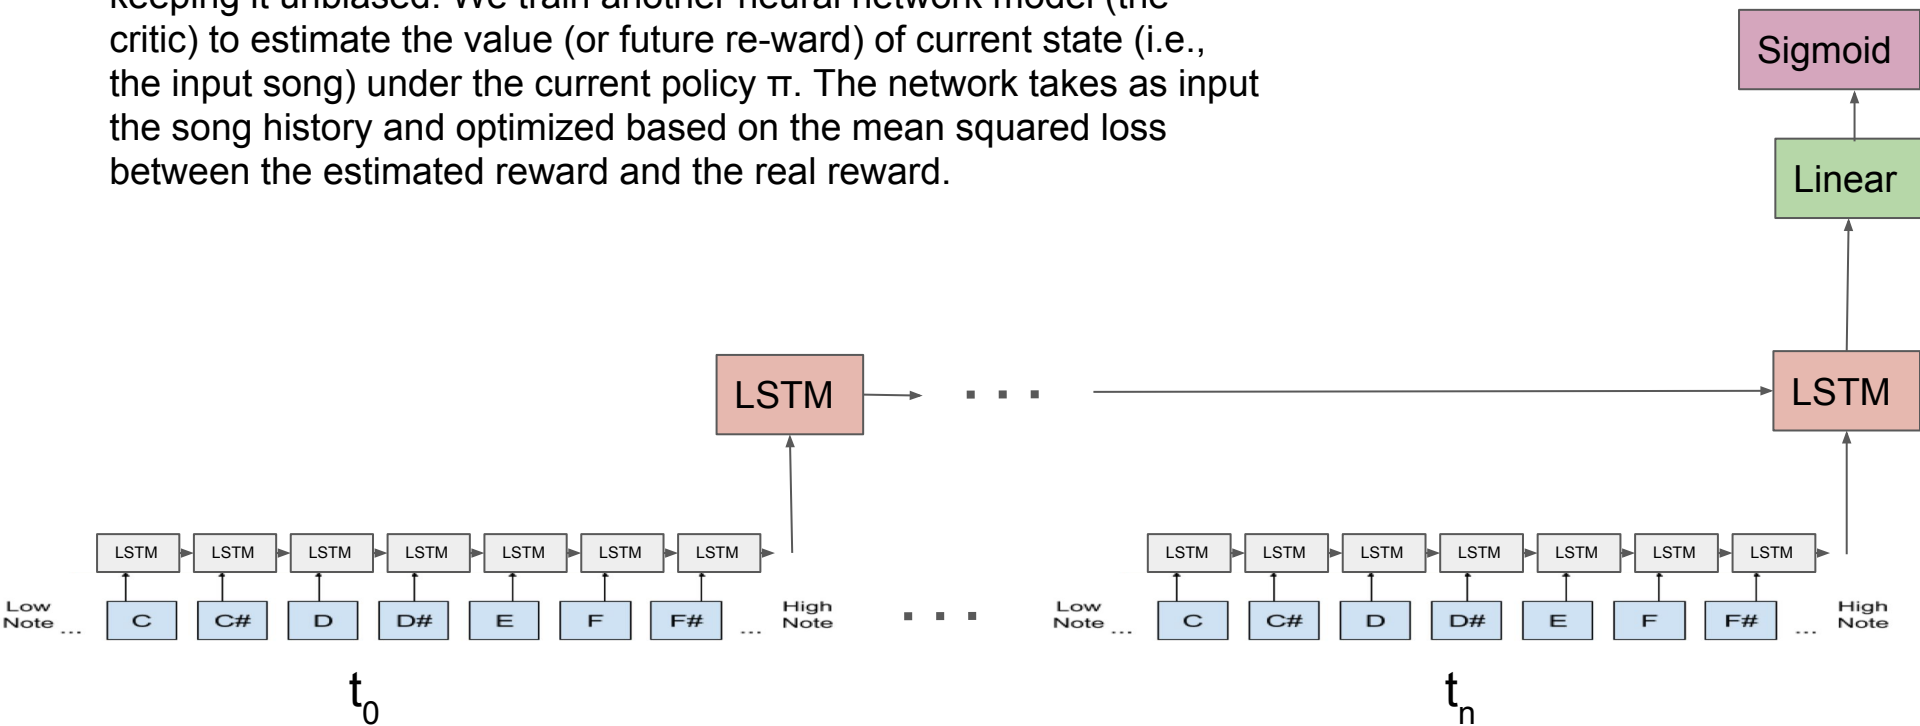
1. Update Discriminator n times
 - a. Sample (X, Y) from real data
 - b. Sample $Y_G \sim G(\cdot|X)$ from generator
 - c. Update D using (X, Y) as positive examples and (X, Y_G) as negative examples.
2. Update Generator m times
 - a. Sample (X, Y) from real data
 - b. Sample $Y_G \sim G(\cdot|X)$ from real data
 - c. Compute Reward r for (X, Y_G) using D.
 - d. Update G on (X, Y_G) using reward r

Policy Gradient of generator:

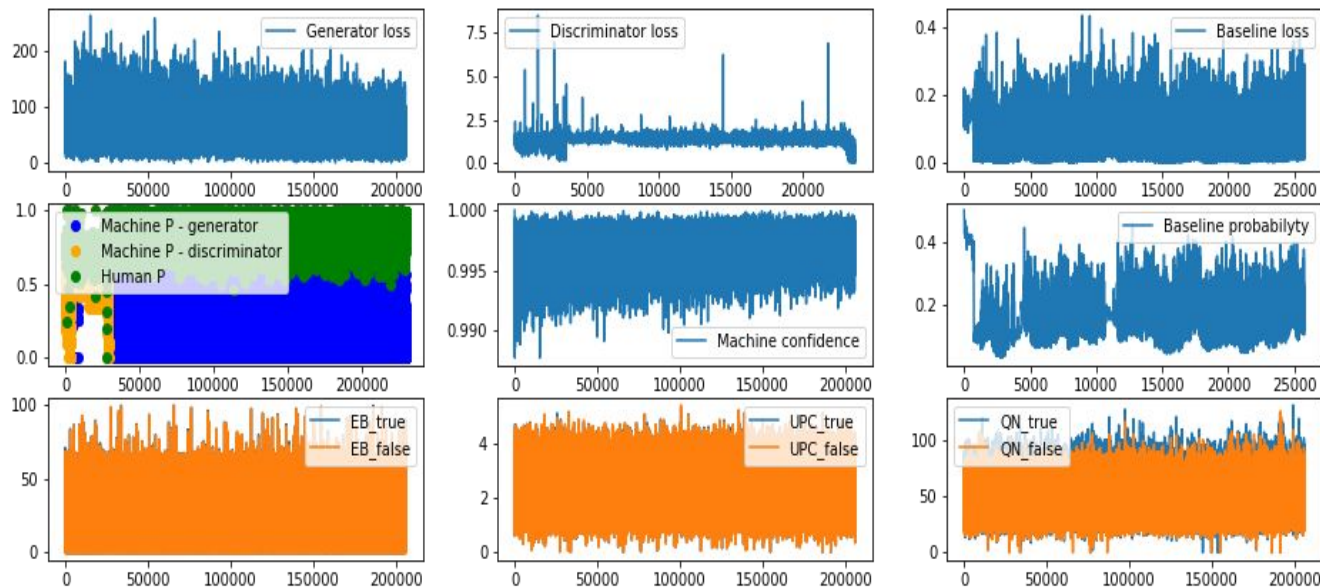
$$\begin{aligned}\nabla J(\theta) &\approx [Q_+(\{x, y\}) - b(\{x, y\})] \\ &\quad \nabla \log \pi(y|x) \\ &= [Q_+(\{x, y\}) - b(\{x, y\})] \\ &\quad \nabla \sum_t \log p(y_t|x, y_{1:t-1})\end{aligned}$$

Baseline

Baseline value use to reduce the variance of the estimate while keeping it unbiased. We train another neural network model (the critic) to estimate the value (or future re-ward) of current state (i.e., the input song) under the current policy π . The network takes as input the song history and optimized based on the mean squared loss between the estimated reward and the real reward.



Typical scenario of learning



EB - ratio of empty bars (in %).

UPC - number of used pitch classes per bar (from 0 to 11)

QN - ratio of "qualified" notes (in %)

Loss of Discriminator:

$$Loss_d = -\log(1 - Q_{+_{false}}) - \log(Q_{+_{true}})$$

Loss of generator:

$$Loss_g = RL_{loss} + Classic_{loss}$$

Loss of Baseline:

$$Loss_b = \sum_i (Q_+(x, y_i) - b(x))^2$$

EB, UPC, QN metrics were proposed in *MuseGAN: Multi-track Sequential Generative Adversarial Networks for Symbolic Music Generation and Accompaniment (2018)*

Learning technique. Comments on Methods

- Teacher Forcing - Use true sample as generator prediction
 - Helps to reach EB, UPC, QN metrics
- Monte Carlo - Sample negative examples several times from generator
 - Does not help much
- Baseline Value
 - Speeds up generator training
- Pretrain generator
- Pretrain discriminator
- Use noise as negative sample to pretrain discriminator
 - Speed up discriminator training
- Use classical cross entropy loss along with RL loss for generator
 - Generated music stay harmonic after a long training



References

1. Generator Architecture:
 - a. DeepJ: Style-Specific Music Generation (2018 [arXiv:1801.00887](#))
 - b. Generating Polyphonic Music Using Tied Parallel Networks (2016).
 - c. Attention Is All You Need (2017 [arXiv:1706.03762](#))
2. Discriminator Architecture:
 - a. A Hierarchical Neural Autoencoder for Paragraphs and Documents (2015 [arXiv:1506.01057v2](#))
3. Learning technique:
 - a. Adversarial Learning for Neural Dialogue Generation ([arXiv:1701.06547](#))
 - b. Simple statistical gradient-following algorithms for connectionist reinforcement learning (1992)
 - c. Image-to-Image Translation with Conditional Adversarial Networks (2017 [arXiv:1611.07004v2](#))
 - d. SeqGAN: Sequence Generative Adversarial Nets with Policy Gradient ([arXiv:1609.05473](#))
4. Metrics:
 - a. MuseGAN: Multi-track Sequential Generative Adversarial Networks for Symbolic Music Generation and Accompaniment ([arXiv:1709.06298](#))

Thank you for your attention!

Ivan.Anokhin@skoltech.ru

Egor.Nuzhin@skoltech.ru

